

PETER CHEN

lc3826@columbia.edu
peterlauukchen.github.io

RESEARCH STATEMENT

My current focus is LLM reasoning and alignment via optimization theory and reinforcement learning design. My areas of expertise include:

1. **LLM Post Training:** Reasoning via RLVR; Alignment (RLHF/DPO); Agent Training.
2. **Optimization/RL Theory:** Optimal Transport; Non-convex Optimization; Minimax Optimization; Zeroth-order Optimization; Multi-agent Reinforcement Learning.

EDUCATION

Columbia College, Columbia University New York, NY
B.A in Mathematics, Computer Science
May 2026

Advisor: Andrew Blumberg (Math), Tianyi Lin (IEOR)

EXPERIENCE

Research Intern | AI Lab, Princeton University Princeton, NJ
Hosted by Prof. Mengdi Wang Feb 2025 – Dec 2025

Topic: LLM RL Reasoning; LLM Agent Training

Research Intern | Institute of Data Science, HKU Hong Kong
Hosted by Prof. Yue Xie, Prof. Qingpeng Zhang May 2024 – Aug 2024

Topic: Neural Optimal Transport, Convex Networks

Teaching Assistant | Department of Mathematics, Columbia University

TA for MATH 2500 Analysis & Optimization over SP24, FA24, SP25, FA25, SP26

PUBLICATIONS

[1] ComPO: Preference Alignment via Comparison Oracles

Peter Chen, Xi Chen, Wotao Yin, Tianyi Lin

Advances in Neural Information Processing Systems 38 (NeurIPS 2025)

[2] Exploration v.s. Exploitation: Rethinking RLVR through Clipping, Entropy, and Spurious Reward

Peter Chen, Xiaopeng Li, Ziniu Li, Xi Chen, Wotao Yin, Tianyi Lin

Proceedings of the International Conference on Learning Representations (ICLR 2026)

[3] Reward-free Alignment for Conflicted Objectives

Peter Chen, Xiaopeng Li, Xi Chen, Tianyi Lin

Under Review, ICML 2026

[4] GenEnv: Difficulty-Aligned Co-Evolution Between LLM Agents and Environment Simulators

Jiacheng Guo, Ling Yang*, Peter Chen*, Qixin Xiao*, Yinjie Wang, Xinzhe Juan, Jiahao Qiu, Ke Shen, Mengdi Wang*

arxiv-2512.19682

[5] Stepwise Guided Policy Optimization: Coloring your Incorrect Reasoning in GRPO

Peter Chen, Xiaopeng Li, Ziniu Li, Xi Chen, Tianyi Lin

Transactions on Machine Learning Research, Under Review

[6] Displacement-Sparse Neural Optimal Transport

Peter Chen, Yue Xie, Qingpeng Zhang

arxiv-2502.01889

[7] 3D Cell Oversegmentation Correction via Geo-Wasserstein Divergence

Peter Chen, Bryan Chang, Olivia Creasey, Julie Sneddon, Zev Gartner, Yining Liu

Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV 2026)

[8] SICNN: Sparsity-induced Input Convex Neural Network for Optimal Transport

Peter Chen, Yue Xie, Qingpeng Zhang

NeurIPS 2024 Optimization for Machine Learning

TALKS

2025 INFORMS Annual Meeting, Atlanta

Oct 2025

Invited Speaker; *LLM Post Training: Turning “Trash” Samples into Value*

SERVICES

Reviewers for: *Conference on Neural Information Processing Systems (NeurIPS), International Conference on Learning Representaton (ICLR), International Conference on Machine Learning (ICML), AAAI Conference on Aritificial Intelligence (AAAI), Conference on Parsimony and Learning (CPAL), Transactions on Machine Learning Research (TMLR)*